



Coordination Languages: Tuple-based and Event-based

Franco Zambonelli
Ottobre 2018



Problems of Distributed Programming

- In programming distributed applications:
 - Often, there are not strict pre-defined roles for the components → processes are peers from the viewpoint of their role
 - Encapsulation → would be great to enable interactions independently of the characteristics of processes
 - We need a mechanism for explicitly designing and controlling interactions between the components → rules for coordination
 - Need to minimize coupling between components → especially needed for dynamic environment (e.g., mobile computing and IoT) where processes might not know a priori each other



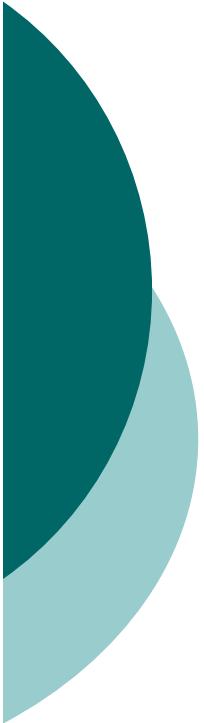
Distributed Object-based Programming

Limitations

- Based on client-server interactions with clear distinction in roles between processes
- The process structure - the number of processes and their relationships - determines the program structure → no explicit design of interactions -> no explicit design of the coordination rules/constraints
- Processes (objects) must know each other and their location and characteristics

Therefore

- Limited openness
- Hard to deal with dynamic environment
- Conceptually not scalable



The way out - Coordination

“Coordination is managing dependencies between activities.”

(From Coordination Theory)

“Coordination is the process of building programs by gluing together active pieces.”

(Carrier and Gelernter, Linda creators)

Coordination includes communication + synchronization



Concurrent/Distributed Programming = Computation + Coordination

(Carriero and Gelernter, Linda
creators)



Coordination: Models, Languages, & Systems

- The **model** is a Triple (**E, M, L**)
 - E** - Coordinables, or **entities** being coordinated (agents, processes, tuples)
 - M** - Coordination media used to coordinate the entities (channels, shared variables, tuple spaces)
 - L** - Coordination **laws** → semantic framework of the model (guards, synchr. constraints)
- **Coordination languages**
 - “The Linguistic embodiment of a coordination model”
- **Coordination system**
 - A middleware to support the model



Coordination Languages must Support:

- Separation between computation and coordination.
- Distribution and (if needed) transparency
 - Also decentralization
- Dynamics (no need to link the activities and the roles before runtime) and Adaptivity
- Interoperability → Openness
- Encapsulation (implementation details are hidden from other components)



Leading Coordination Models and Languages

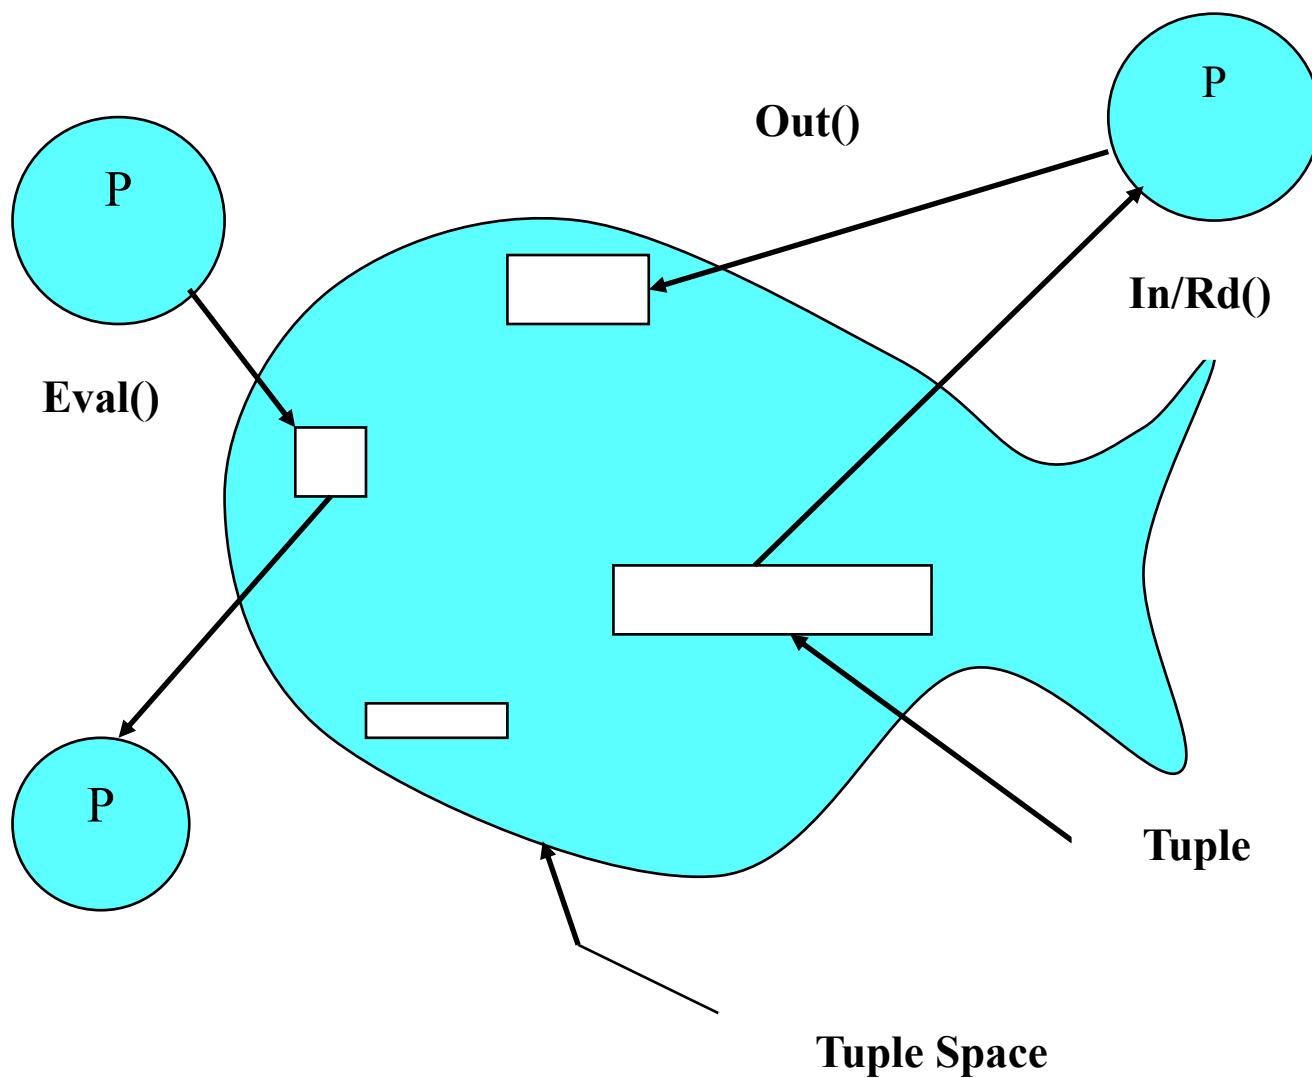
- Tuple Space Model (coordination through Shared Data Space)
 - Linda, TSpaces, JavaSpaces, Jini
- Event-based Models
 - We will see it as an extension to Tuple Spaces
- Multiagent Protocols
 - See Lesson on Multiagent Systems



Tuple spaces and Linda

- Based on *generative* communication:
 - Processes communicate by inserting or retrieving data objects (called Tuples) from shared Data Space (called Tuple space)
 - The tuple space is a multiset of “bag of tuples” – unstructured DB
- Interoperability is achieved: Linda places no constraints on the underlying language or platform for the processes

Tuple Space Concept





Tuples

- A tuple is a structured series of typed fields.
 - Example:
 - (“**a string**”, **15.01**, **17**)
 - **Int index = 5; float value = 4.45;**
 - (“**Vector**”, **index**, **value**)
- Depending on the specific implementation we can have OO tuples or logic tuples, e.g., OO
- (**String “hello”**, **Counter c**, **Person f**)



Operations

out (t) insert a tuple t into the Tuple space (non-blocking)

in(at) if there is a tuple **t** from the tuple space matching the template tuple **at** (often called “antituple”), it finds and removes it; **block** until a matching tuple is found – a template tuple is a normal tuple that can have some “formal” variable **?v** , i.e., fields undefined in value. These will assume the values found in the marching tuple, i.e., like **v = value of corresponding field in t**

rd(at) like in(at) but the t not removed blocking

Bool inp(at) tests the existence of a tuple without blocking

eval(t) add an active tuple t to the tuple space – basically it is the creation of a new process (not often used in real implementations, where process are created as usual)

Tuple matching

Let **at(i)** denote the i_{th} field in the (template) tuple **at**.

A tuple **at** given in a **in(at)** or **rd(at)** operation “matches” a tuple **t** in the tuple space iff:

1. **t** and **at** have the same number of fields, and
2. for each field
 - if **at(i)** is a value then **at(i) = t(i)**

or

if **at(i)** is of the form **?x** then **t(i)** is a valid value for the type of variable **x**

If more than one tuple in the tuple space matches, then one is selected **nondeterministically**.

As a result of tuple matching if **at(i)** is of the form **?x**, then **x := t(i)**



Examples of Tuple matching

N-element vector stored as n tuples in the Tuple space:

```
(“v”, 1, FirstElt)  
(“v”, 2, SecondElt)  
...  
(“v”, n, NthElt)
```

To read the jth element and assign it to x, use

```
rd (“v”, j, ?x )
```

To change the ith element, use

```
in (“v”, i, ?OldVal)  
out (“v”, I, NewVal);  
/* Note: impossible updating a tuple – need to remove and  
re-insert!!!! */
```



Programming Concurrent and Distributed Applications in Linda

- With Linda
 - We can replicate most of the functionalities of message-oriented and client server systems
- However
 - We can use the tuple space as a more general means to realize dynamic and open forms of interaction

Example 1: Sending and Receiving Messages

- We can use the tuple space to simulate a communication channel between two processes

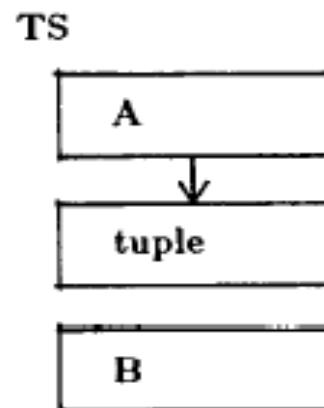


Fig. 1. To send to B, A generates a tuple...

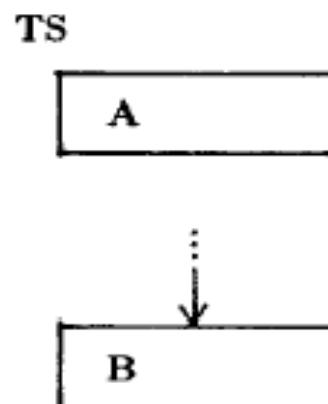


Fig. 2. ... and B withdraws it.

Sending and Receiving messages (Cont.)

```
// Tuple structure (From, To, Index, Data)

Process p
{ int Index = 0; Process q = q_id; Anytype Data = ...;
  out(p,q,Index,Data)
  Index++;
  // further messages...
}

Process q
{ Process sender, p = p_id; Anytype Data;
  in(p, q, ?Index,?Data)}
  // if interested only in messages from a specific p

  in(?sender, q,?Index,?Data)}
  // if interested in messages from anyone
}
```



Sending and Receiving messages (Cont.)

- Message-passing model contains three separate parts:
 1. Targets (to-from)
 2. Tag (message Id or index)
 3. Communication
- In Linda, Tuple is a set of fields, any of which may be a tag (message Id) or target (address) or communication (data).

Example 2: Client-Server

```
client()
{ int index;
...
in ("server index",?index);
out ("server index", index +1);
...
out ("server", index, request);
in ("response", index, ?response);
...
}
```

The index is used to
Associate the client
With its answer

```
server()
{ int index = 1;
out("server index", index);

while (1){
    in ("request", ?index, ? req)
    ...
    out ("response", index++, response);
}
}
```



Example 3: Semaphores

- Semaphores

Initialize: out("sem")

wait or P-operation: in("sem")

signal or V-operation: out("sem")

For counting semaphore execute
out("sem") n times

Example 4: Concurrent Processes

Multiple processes can execute on the same task, splitted in several subtasks, by simply getting the tuple representing the task No need to know a priori how many processes

```
// create task: the element of the matrixes in the tuple space
for (i,j = 1...N)
{  Out(i,j, "A", A[i][j]);
   Out(i,j, "B", B[i][j]);

   In(i,j, "C", ?C[i][j]); // read results
}

// then launch any number of concurrent processes doing this
While(inp(?i,?j, "A", ?int)
{  in(?i,?j, "A", ?A[i][j]);
   in(?i,?j, "B", ?B[i][j]);
  in(?i,?j, "C", A[i][j]+B[i][j]); }
```



Concurrent Processes: Masters and Workers

- Is general:
 - A master process divides work into discrete tasks and puts them into global space (tuple space)
 - Workers retrieve tasks, execute, and put results back into global space
 - Workers notified of work completion by having met some condition → e.g. no more tuples representing task or insertion by the master of a specific tuples
 - Master gathers results from global space

Masters and Workers (Cont.)

```
master() {  
    for all tasks {  
        .../* build task structure*/  
        out ("task", task_structure); }  
    for all tasks {  
        in ("result", ? &task_id, ?  
            &result_structure); }  
}  
  
worker() {  
    while (in ("task", ? &task_structure) {  
        .../* exec task */  
        out ("result", task_id,  
            result_structure); }  
}
```



Example 6: Barrier Synchronization

Each process within some group must wait at a “barrier” until all processes in the group have reached the barrier; then all can proceed.

- Set up barrier:
`out ("barrier", n);`

- Each process does the following:
`in("barrier", ? val);
out("barrier", val-1);
rd("barrier", 0).`

Example 7: Dining Philosophers Problem

```
phil(i)
{   int i;

    while (1) {
        think();
        in ("room ticket");
        in ("chopstick", i)
        in ("chopstick", (i+1)% Num);
        eat ();
        out ("chopstick", i);
        out("chopstick", (i+1)%Num);
        out ("room ticket");
    }
}
```

```
initialize()
{ int i;
    for (i=0; i< Num; i++){
        out ("chopstick", i);
        create_process( phil(i) );
        if (i< (Num-1))
            out("room ticket"); }
}
```

Example 8: Analyzing Sensor Data

```
sensor(id)
{
    if (new interesting sample)
        out("temperature", id, value, timestamp);
}

analyzer()
{ int MAX=0, sensor_MAX, sensor_id, new_value, timestamp;

    while (1) {
        in
        ("temperature", ?sensor_id, ?new_value,?timestamp);

        if(new_value > 0)
            MAX = new_value;
            sensor_max = sensor_id;
    }
}
```



Linda achieves:

- “Good” Nondeterminism
- Structured naming. Similar to “select” in relational DB, pattern matching in AI.

```
in(P, i:integer, j:boolean)  
in(P, 2, j:boolean)
```
- **Time uncoupling** (Communication between time-disjoint processes)
- **Space uncoupling** (no need to share a process name space)



Linda: Limitations

- Events
 - How to handle asynchronous computing
 - Need to capture events without need to polling the tuple space
- The language is although intuitive, but minimal
 - Real world implementation adds several features to make it useful in real-world
- Implementation:
 - How is the tuple space to be implemented in the absence of shared memory in a distributed system?
- Security
 - Need to have access rights and confinement in the access to tuple spaces
- Programmability of Coordination Laws
 - It is possible to change the basic matching rules so as to support different schemes (constraints on interactions or specific composition rules)?



Handling Events

- Add a primitive operation to Linda
 - **Sub(at, Handler)** -- and also **unsub(t)**
 - Non blocking
- When a tuple matching at appears in the tuple space, execute the function Handler
 - In this way, a process does not have to “poll” the tuple space
 - The tuple space becomes the fact an “Event-dispatcher” to handle publish/ subscribe interaction model
- Since tuples are persistent (while events can be temporary) we can consider also adding a “timeout” (max time of life) to tuples
 - **Out(t, timeout)**

Example 8: Analyzing Sensor Data

Event-based Approach

```
sensor(id)
{
    while(1, every 10 seconds)
        out(("sensor", id, value, timestamp),10);
    // the last parameter specify the lease time
}
analyzer()
{ int MAX=0, sensor_MAX, sensor_id, new_value, timestamp;

    sub(("sensor", ?sensor_id, ?new_value,?timestamp), f1);
    // the analyzer can do other stuff without polling the
}

void f1(t) // t is the tuple matching the subscription
{if(t.new_value > 0)
    MAX = t.new_value;
    sensor_max = t.sensor_id;
}
```

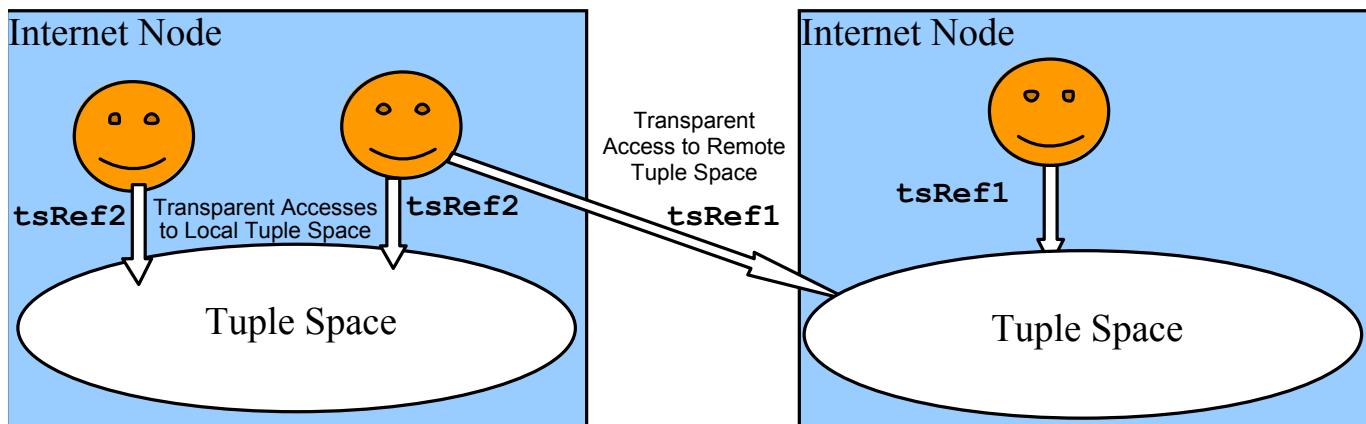


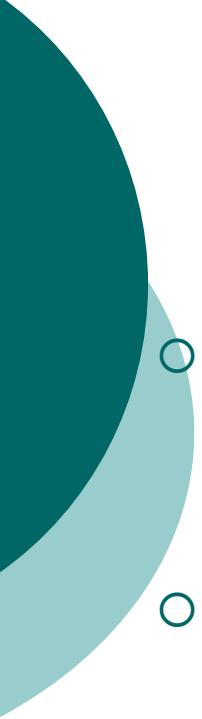
Real-world Example: Java Spaces

- Integrated in J2EE and JINI
- Is the basic engine behind the discovery services of JINI and the JNDI
- Can be used as a stand-alone service to implement tuple-based coordination

JavaSpaces: Coordination Media

- Tuple spaces as (remote) Java objects
- Possible multiple ***independent*** tuple spaces in different nodes
 - tuple spaces identified and accessed via object-references (as in RMI)
 - Accesses can be remote
- References to tuple spaces can be retrieved via the Loopup service of JINI

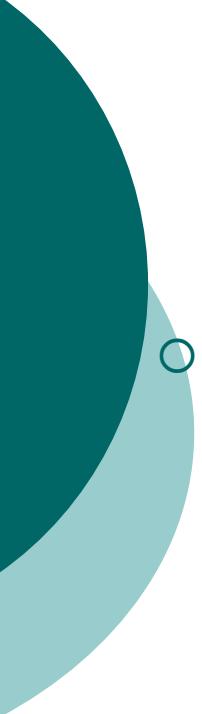




JavaSpaces: Tuples

- Object-oriented Linda operations
 - Tuples are called “Entry”
 - Their fields are objects
- The object tuple *tmp*:
 - subclass of `Entry` ; its instance variable are the tuple fields

```
public class MyTuple extends Entry {  
    public int field1;  
    public float field2; // a field could be any object...  
  
    public MyTuple(int field1, float field2)  
    {this.field1=field1;  
     this.field2=field2}
```



JavaSpaces: Operations

- JavaSpaces interface
 - **Read (Entry tmp, Tr txn, long lease)**
 - **Take (Entry tmp, Tr txn, long lease)**
 - **Write (Entry e, Tr txn, long timeout)**
 - txn can specify a transactions (*all-or-nothing* semantics)
 - lease blocking time for operations (not infinite blocking)
 - timeout time to live for tuples, then they are deleted
- JavaSpaces matching
 - The template must be a class or a superclass
 - Its field with “null” value correspond to ?type of Linda templates



JavaSpaces: Events

- Reactive event notification
- Agents can register in a space as interested in a tuple matching a template, and require to be notified about its insertion

```
notify(template, null, listener, lease.FOREVER, null)
```

- the listener will be sent an event it has to handle when a matching tuple will be inserted in the space
- Can catch a limited set of events!

JavaSpaces: Sample Code (1)

```
import net.jini.core.entry.*;

// definition of the tuple class
public class MessageEntry implements Entry {
    public String content;
    public int value;

    public MessageEntry() {
        } // with no parameters is for initializing templates

    public MessageEntry(String content, int value) {
        this.content = content;
        this.value = value;
    }

    public String toString() {
        return "MessageContent: " + content + " " + value;
    }
}
```

JavaSpaces: Sample Code (2)

```
import net.jini.space.JavaSpace;

public class SpaceClient {
    public static void main(String argv[ ]) {
        try {
            MessageEntry msg = new MessageEntry("Valore", 5);
            // tuple created and initialized

            System.out.println("Searching for a local JavaSpace...");

            Lookup finder = new Lookup(JavaSpace.class);
            JavaSpace space = (JavaSpace) finder.getService();
            // this is the lookup service of JINI

            System.out.println("A JavaSpace has been discovered.");

        // CONTINUE ON NEXT PAGE
    }
}
```

JavaSpaces: Sample Code (3)

```
// CONTINUE FROM NEXT PAGE

System.out.println("Writing a message into the space...");  
space.write(msg, null, 60*60*1000);

MessageEntry template = new MessageEntry();  
template.content = "Valore"  
// a template has some fields null

System.out.println("Reading a message from the space...");

MessageEntry result =  
(MessageEntry) space.read(template, null, Long.MAX_VALUE);

System.out.println("The message read is:"  
+ result.content);
} catch(Exception e) {  
e.printStackTrace();
}}
```



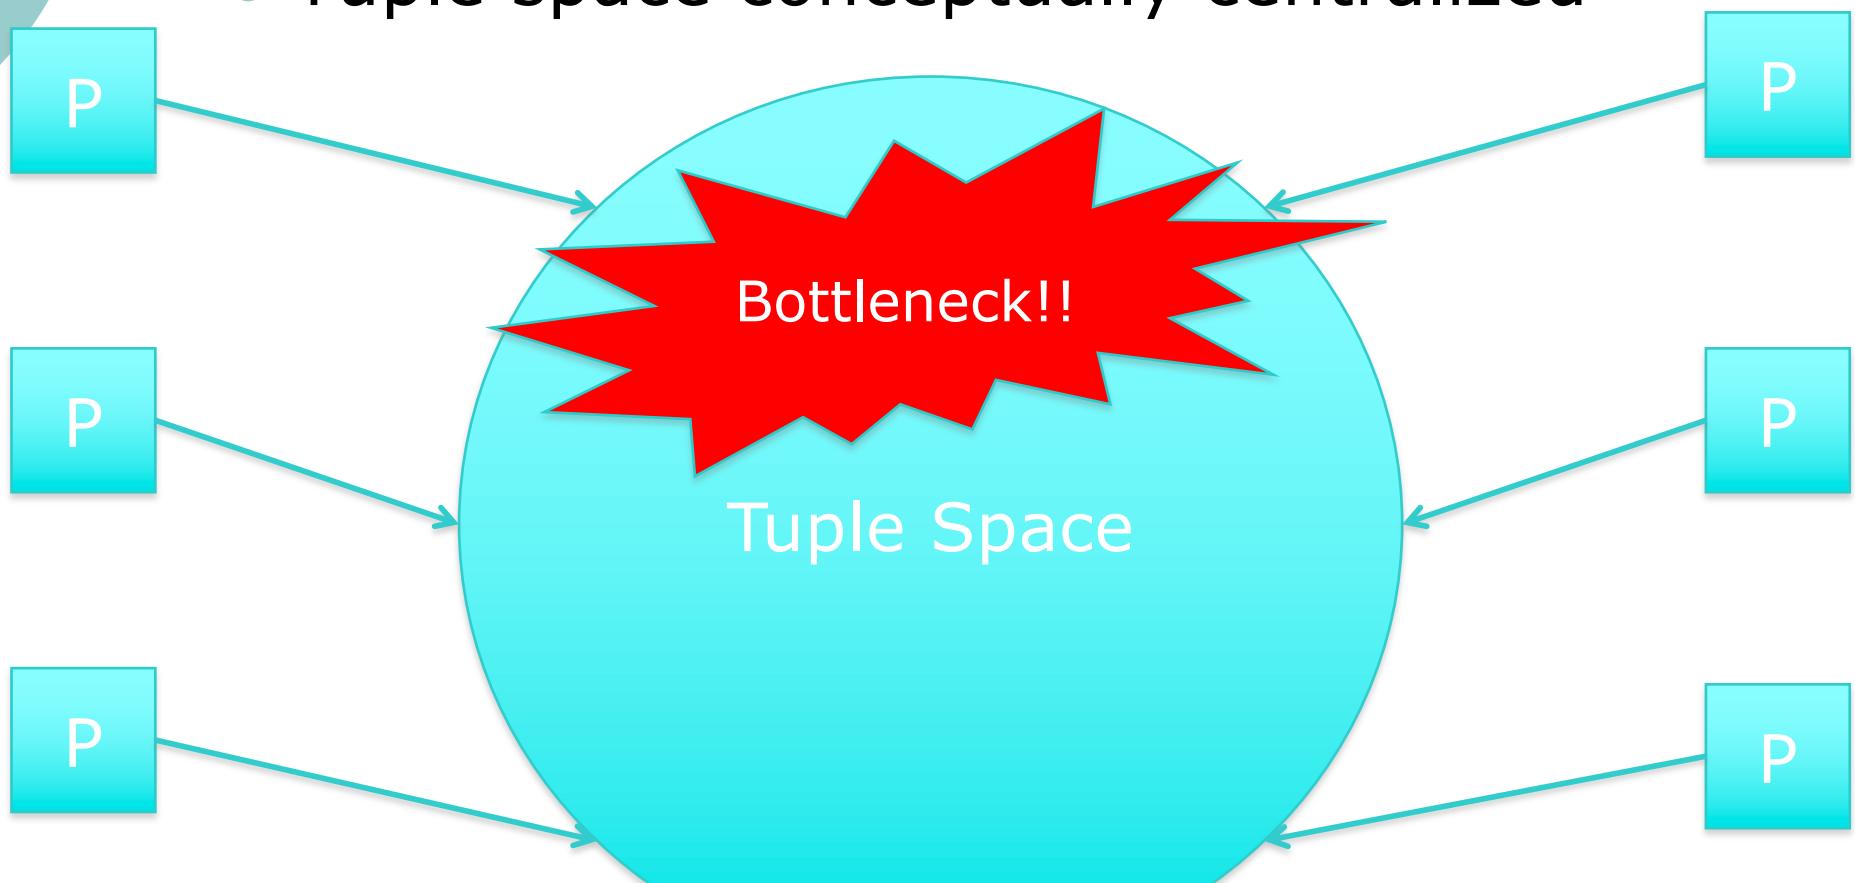
Implementing Tuple Spaces

- Two main approaches
- Global (virtually unique) distributed tuple space
- Local multiple tuple spaces (set of localities)

Global Tuple Space

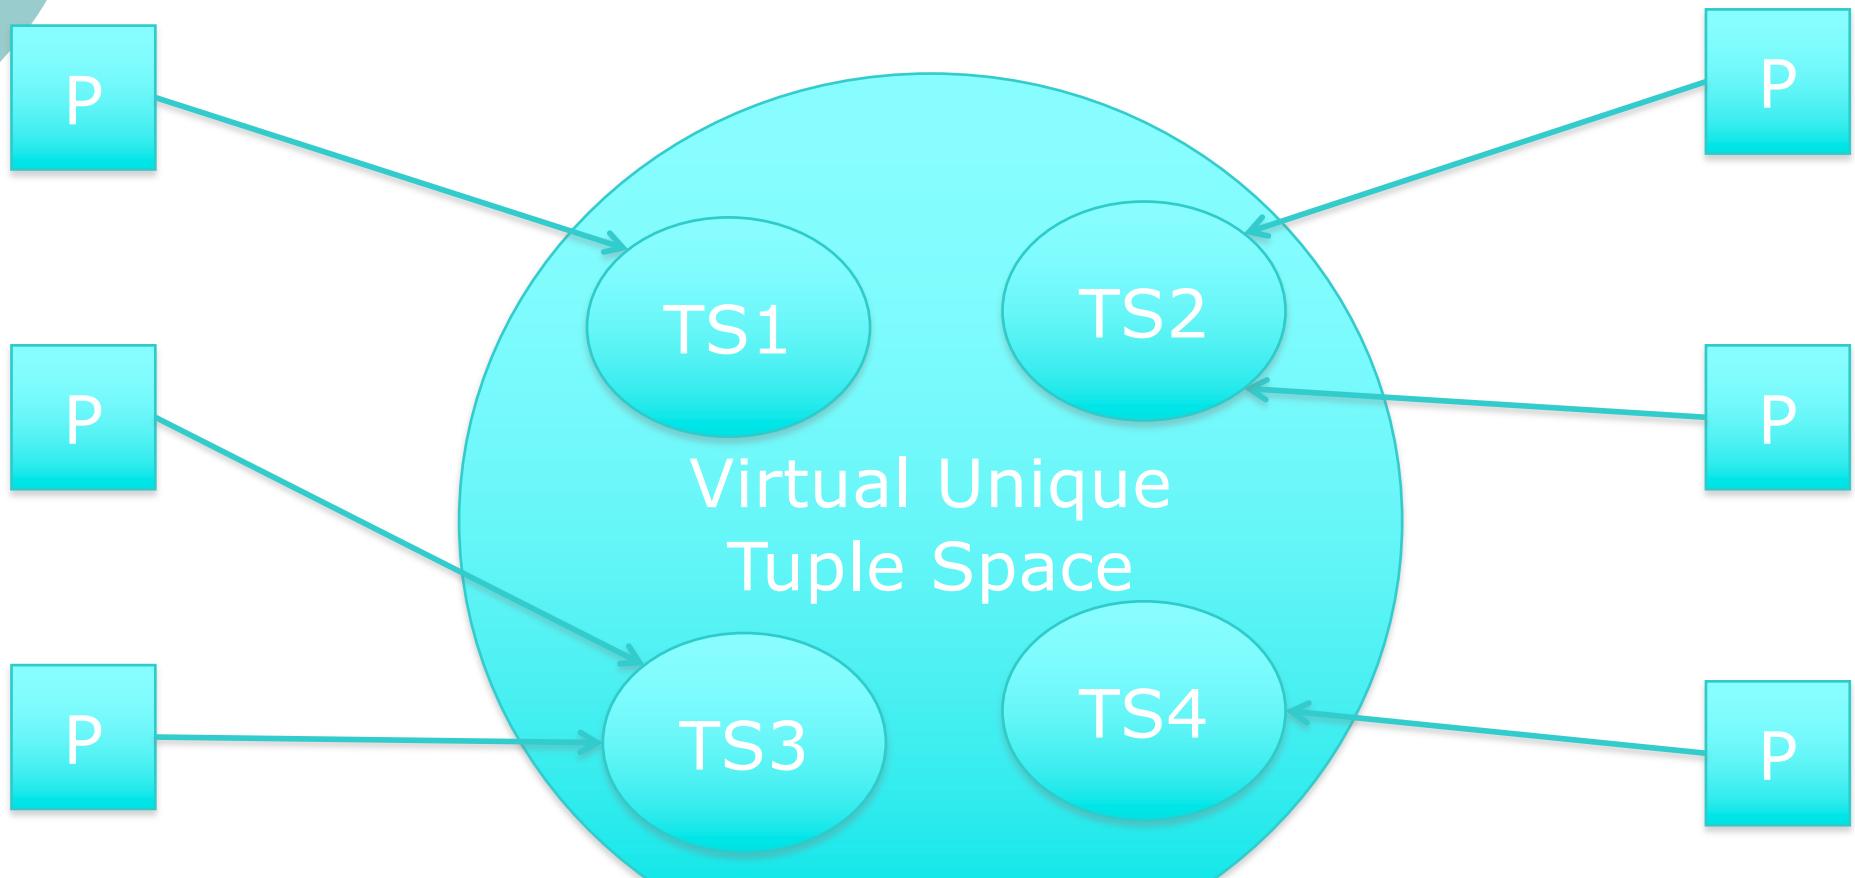
- Dilemma

- Processes distributed
- Tuple space conceptually centralized



Distributed Implementation of a Virtually Unique Tuple Space

- A multiplicity of tuple spaces
- Coordinate to appear as a single one
- Locally accessed by processes





Goal of a Distributed Implementation

- Reachability

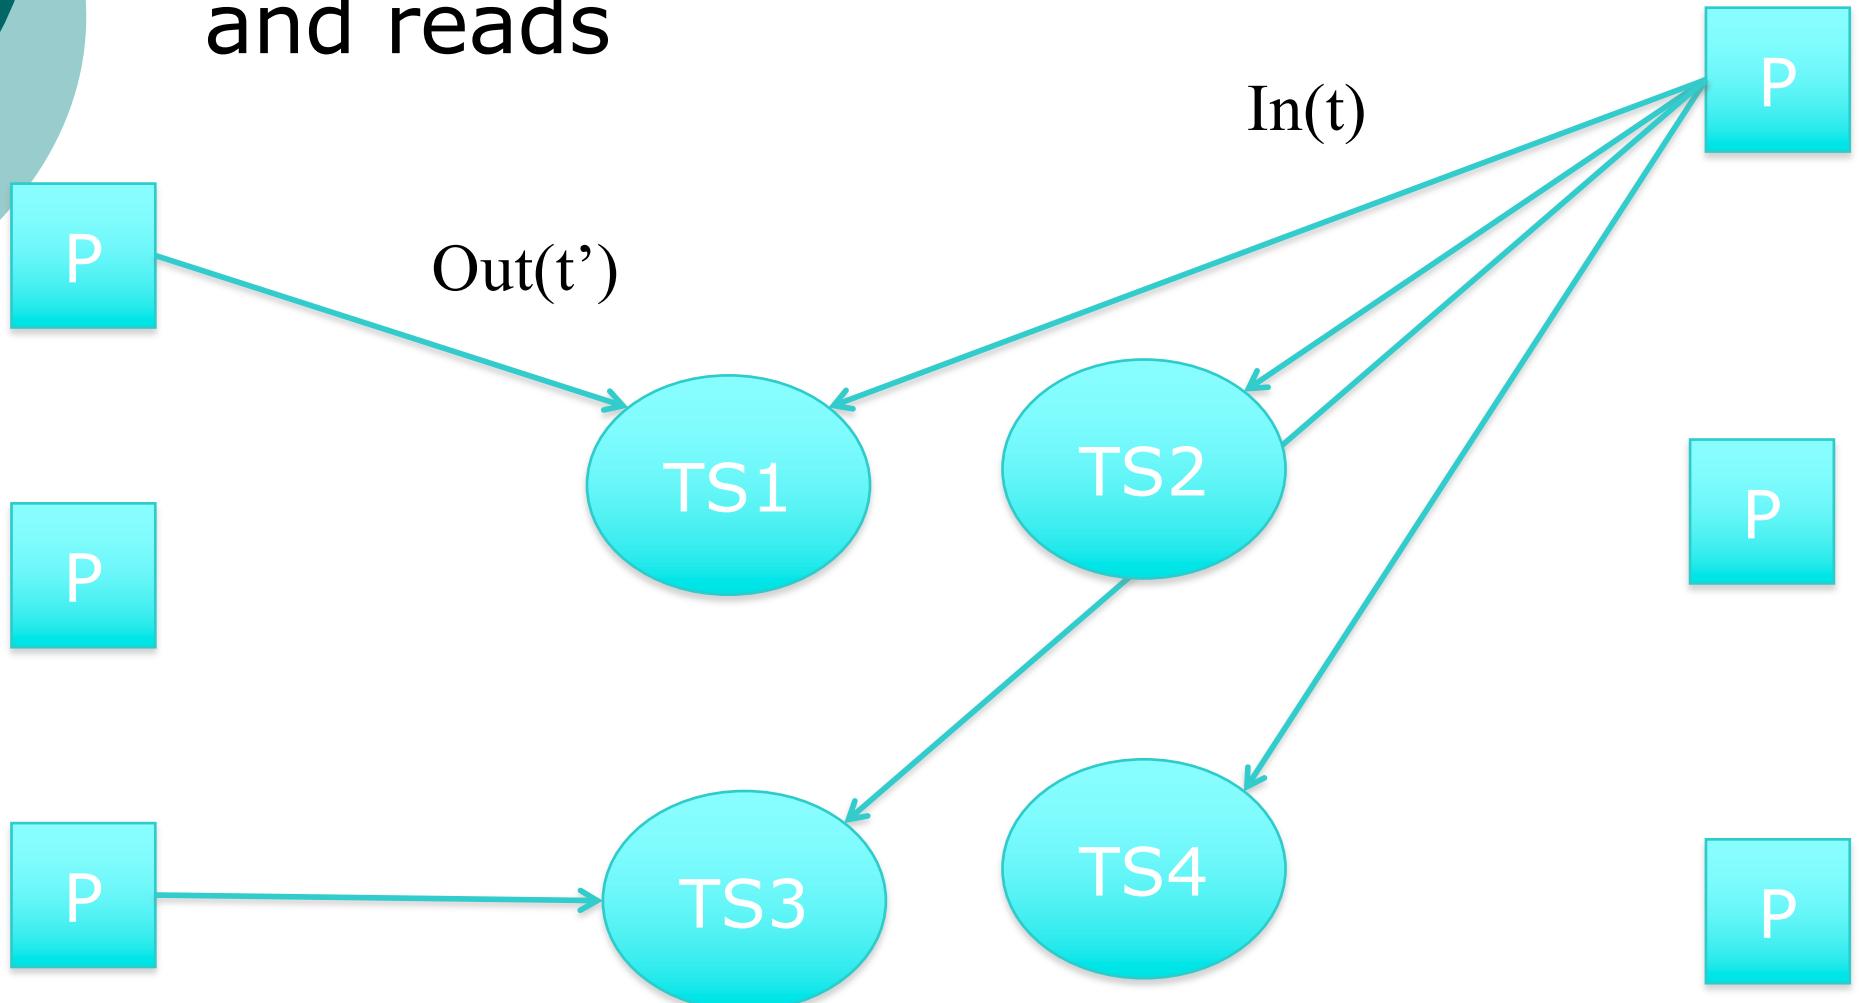
- For any `in(t)` and `read(t)` operations
- If there exists a tuple t' matching t – this must be found

- Consistency

- In case of `ins` operations
 - A. Only a single instance of t' could be returned to different processes (a single tuple cannot be extracted multiple times)
 - B. process could extract a single t' with an operation

Strategies: Replicative Ins

- Local copy of outs, broadcast of ins and reads



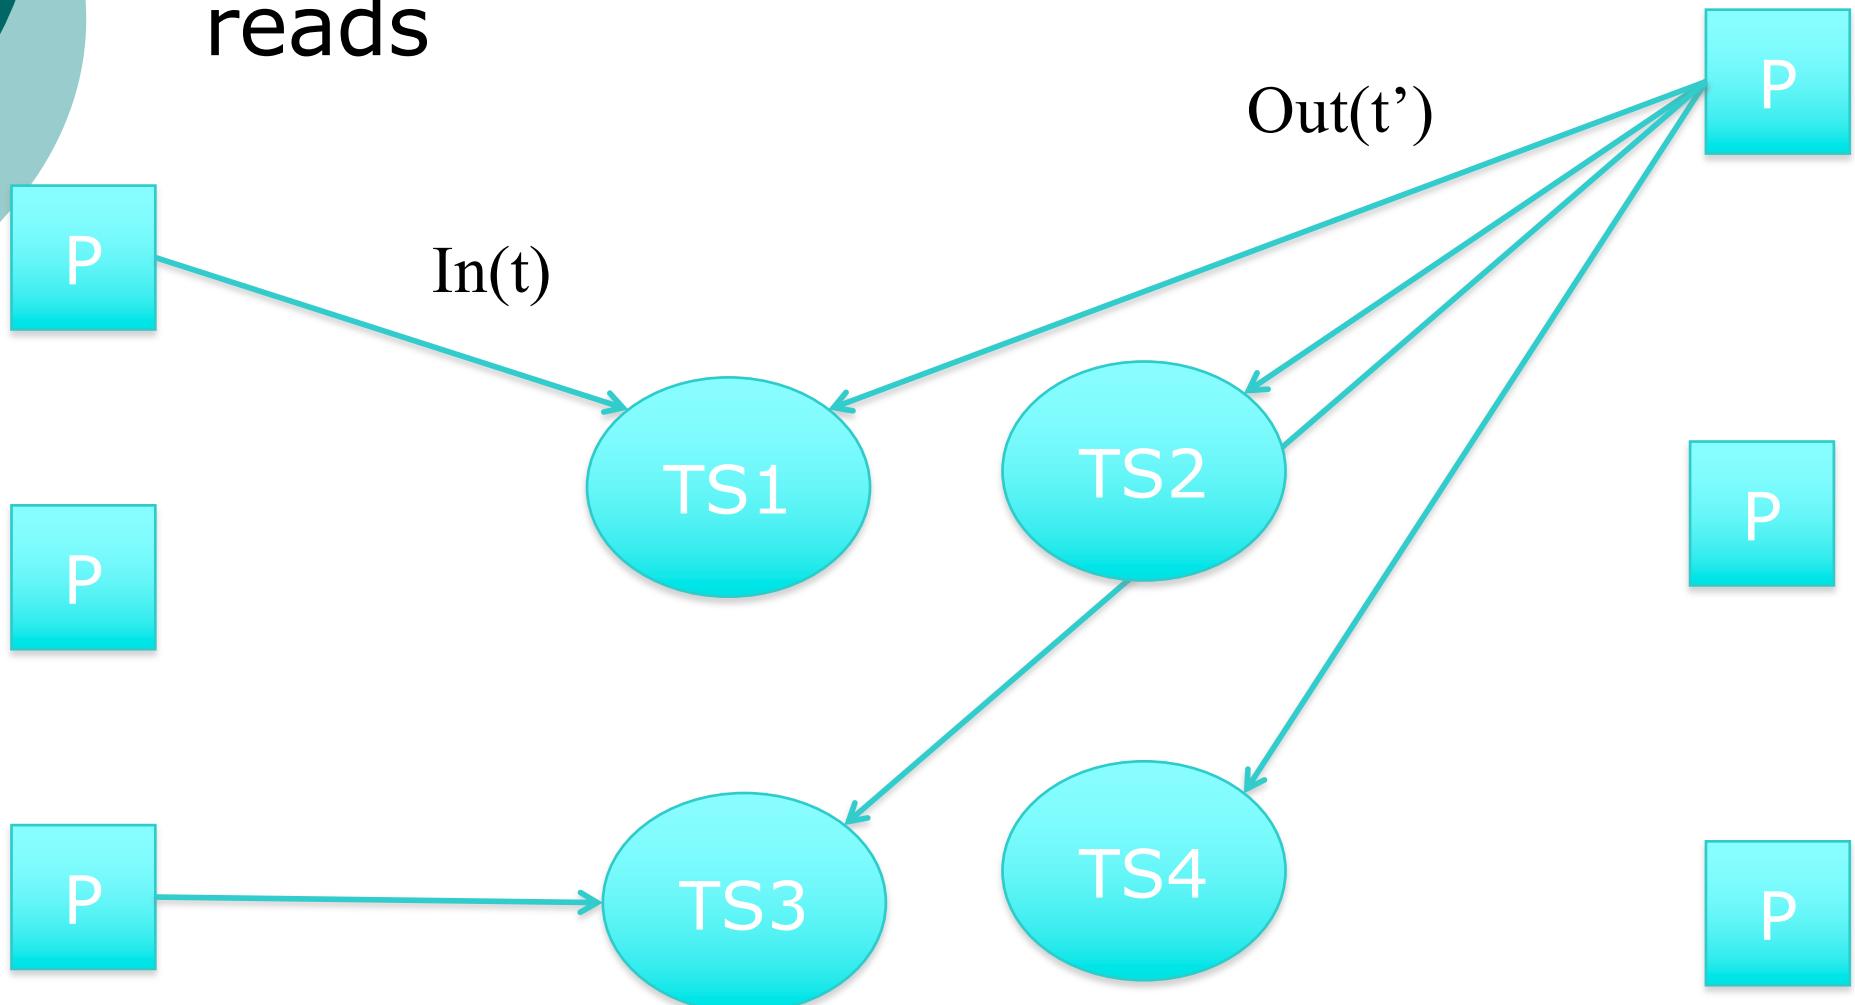


Strategies: Replicative Ins

- Reachability
 - Ensured by construction
- Consistency
 - A. Given that a tuple t' exists only on one node, it is not possible for different processes to extract the same tuple
 - B. if the same process discover multiple matching tuples on different nodes, these are re-routed on such nodes – this is true both both `ins` and `rds`. This implies executing the operation withing a **distributed transaction algorithm**.
- Important: a blocking `in(t)` implies replicating the `in(t)` request on all nodes until it is satisfied, this once a match has been found on a node all replicated `ins` have to be deleted

Strategies: Replicative Outs

- Broadcast of outs, Local of ins and reads





Strategies: Replicative Outs

- Reachability
 - Ensured by construction
- Consistency
 - A. Once the in found a local matching tuple t' , it must coordinate a process for extracting the same tuple t' on all nodes. Only if all replicas of t' have been successfully extracted t' can be delivered to a process. This is a problems of atomic acquisition of multiple resources, similar to the dining philosophers → requires **distributed mutual exclusion algorithms**
 - B. No problem of discovering multiple tuples on different nodes

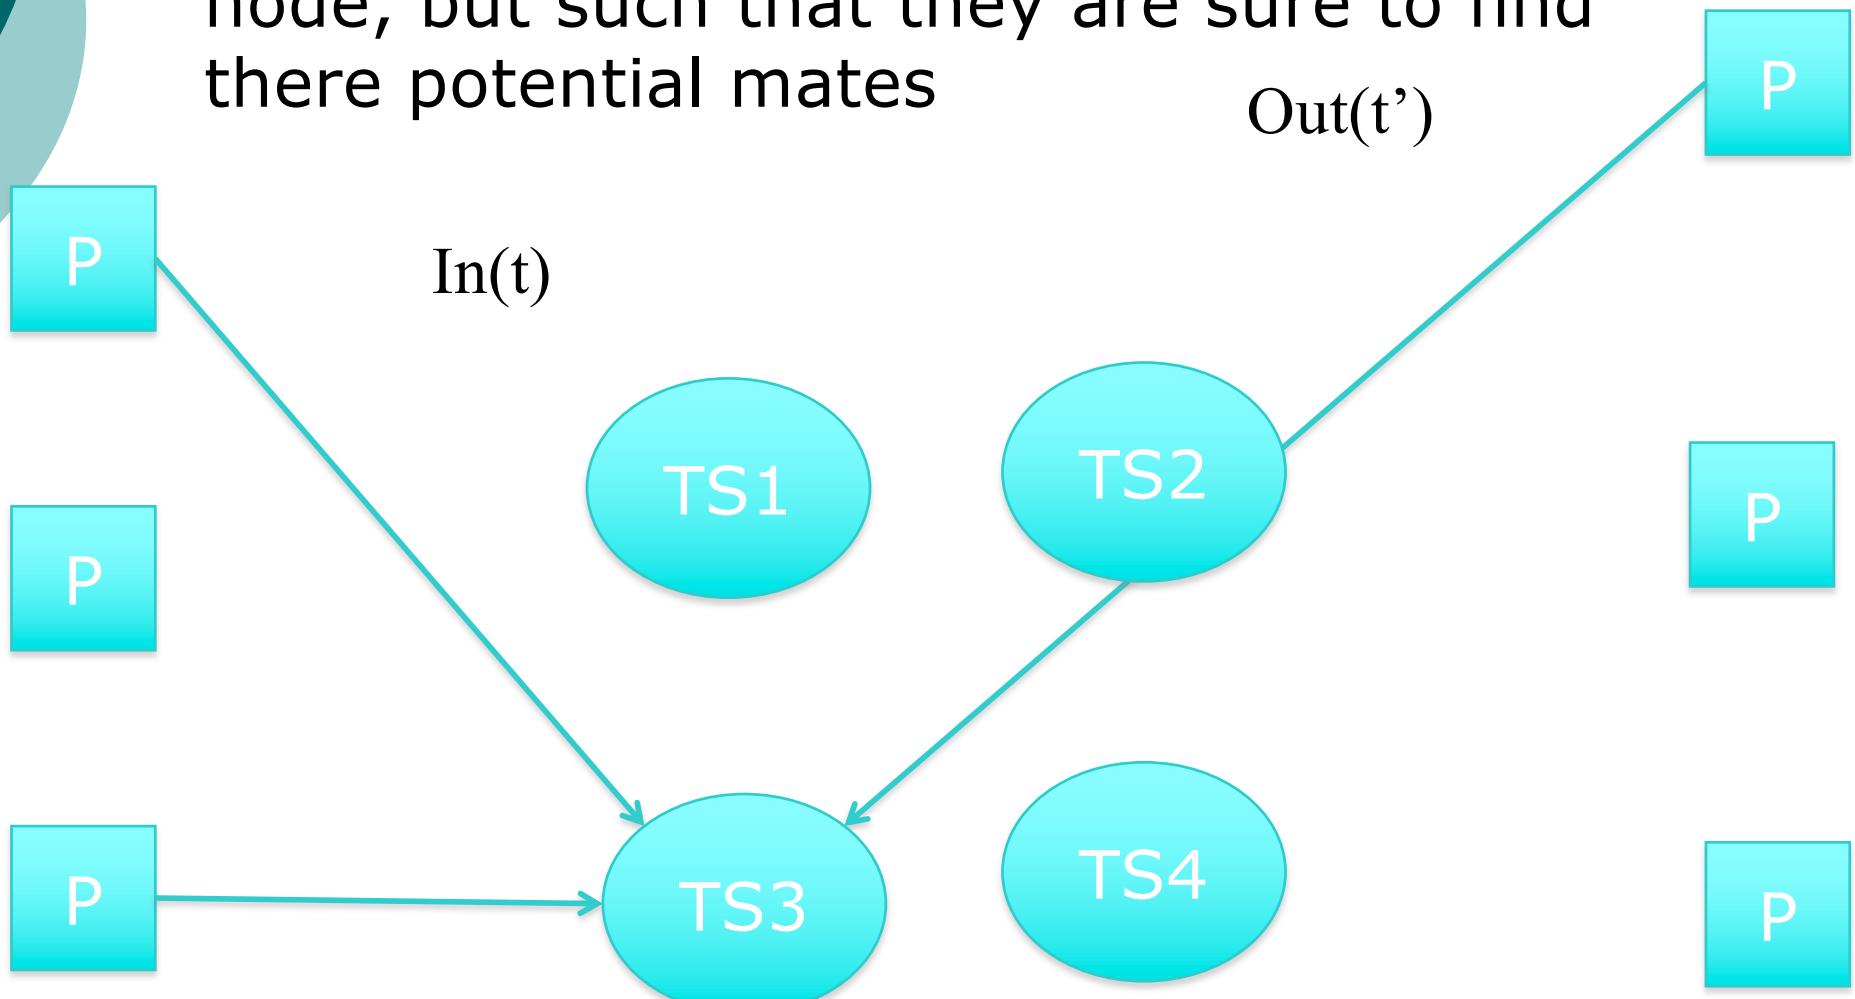


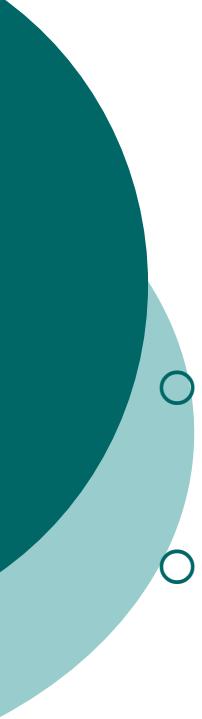
Pros Cons of Replicative Policies

- Pro: Exploit locality in the accesses
 - For either ins/rds or outs a process can access the most local tuple space
- Cons:
 - Broadcasting in large-systems can be expensive
 - The transactions and mutual exclusion algorithms to be implemented can be complex

Strategies: Hash-based

- Assign tuples and templates to a single node, but such that they are sure to find there potential mates





Strategies: Hash-based

- Given N nodes that are implementing the tuple space
- Define a function H that maps a tuple (or a template) into one of the N nodes
 - $H: T \rightarrow 1..N$
- Make sure that, for any tuple t and for any potentially matching template t'
 - $H(t)=H(t')$
- Typically, the hash function applies either on field string characterising all tuples of a given kind, or on the structure of the tuples

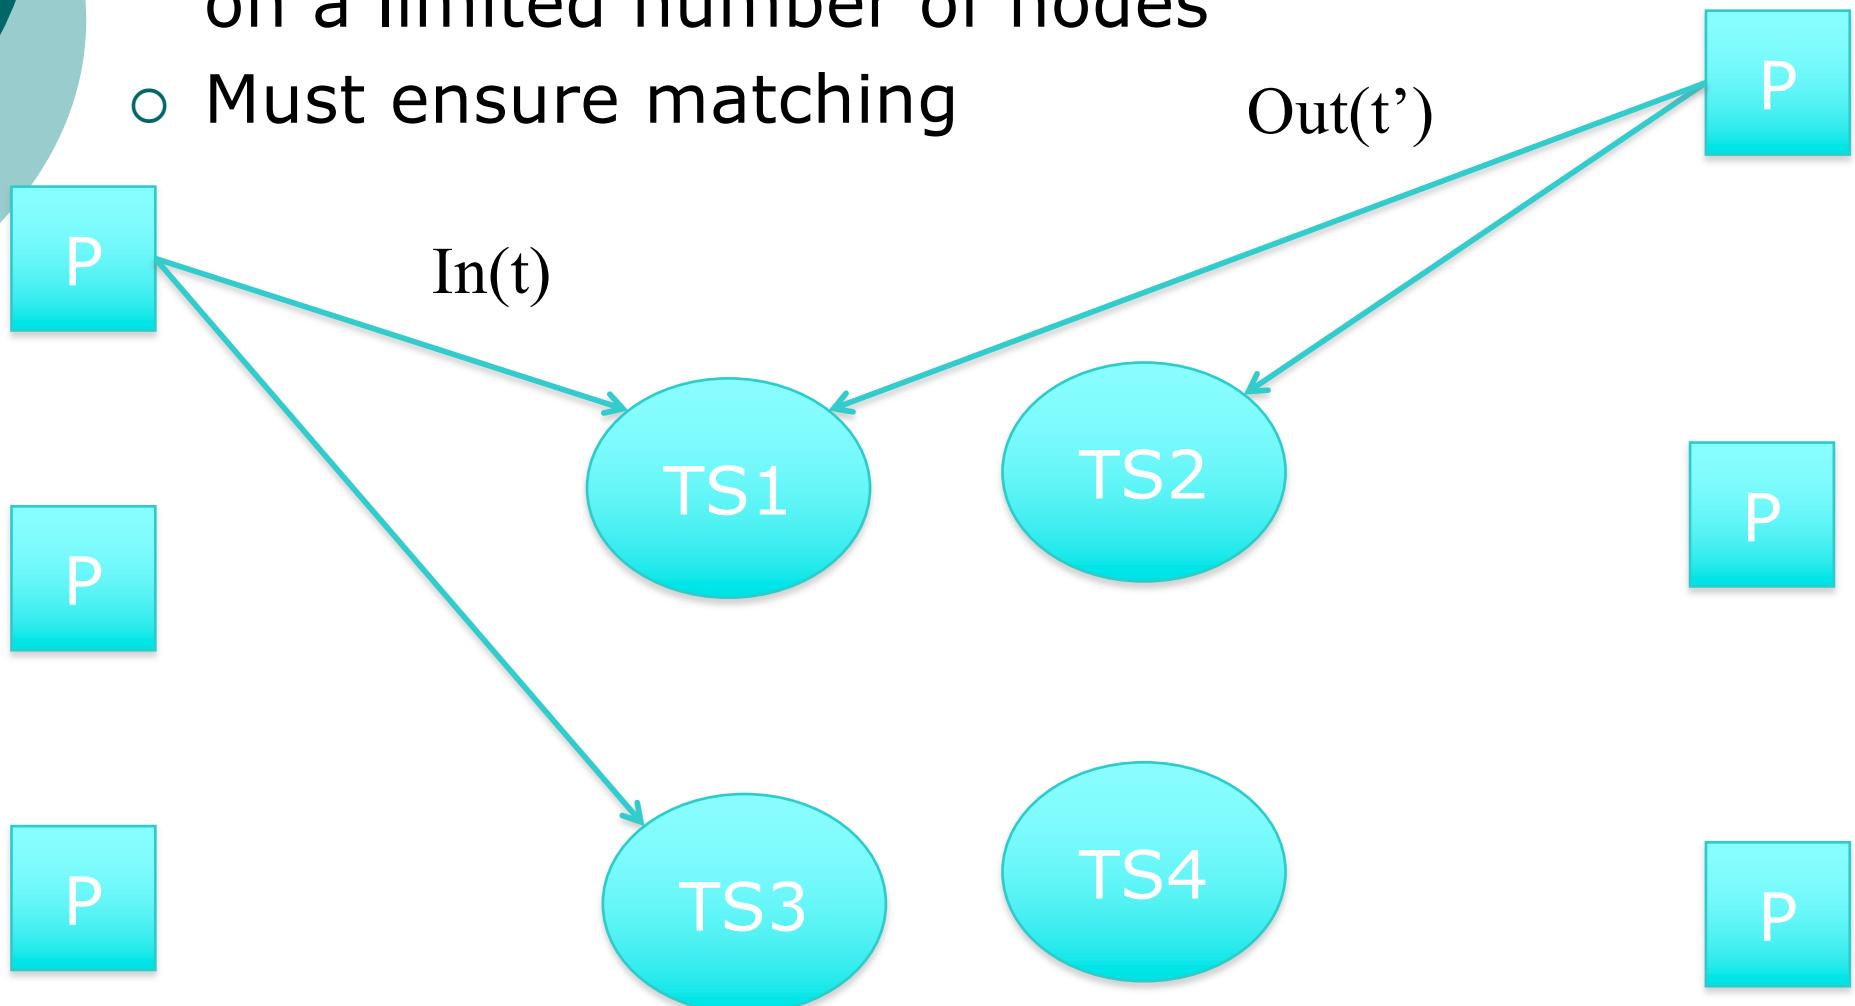


Strategies: Hash-based

- Reachability guaranteed by construction
- Consistency guaranteed because there is no replication
- Problem:
 - Do not exploit locality in the accesses
 - The node that a process accesses is not related to the location of the processes

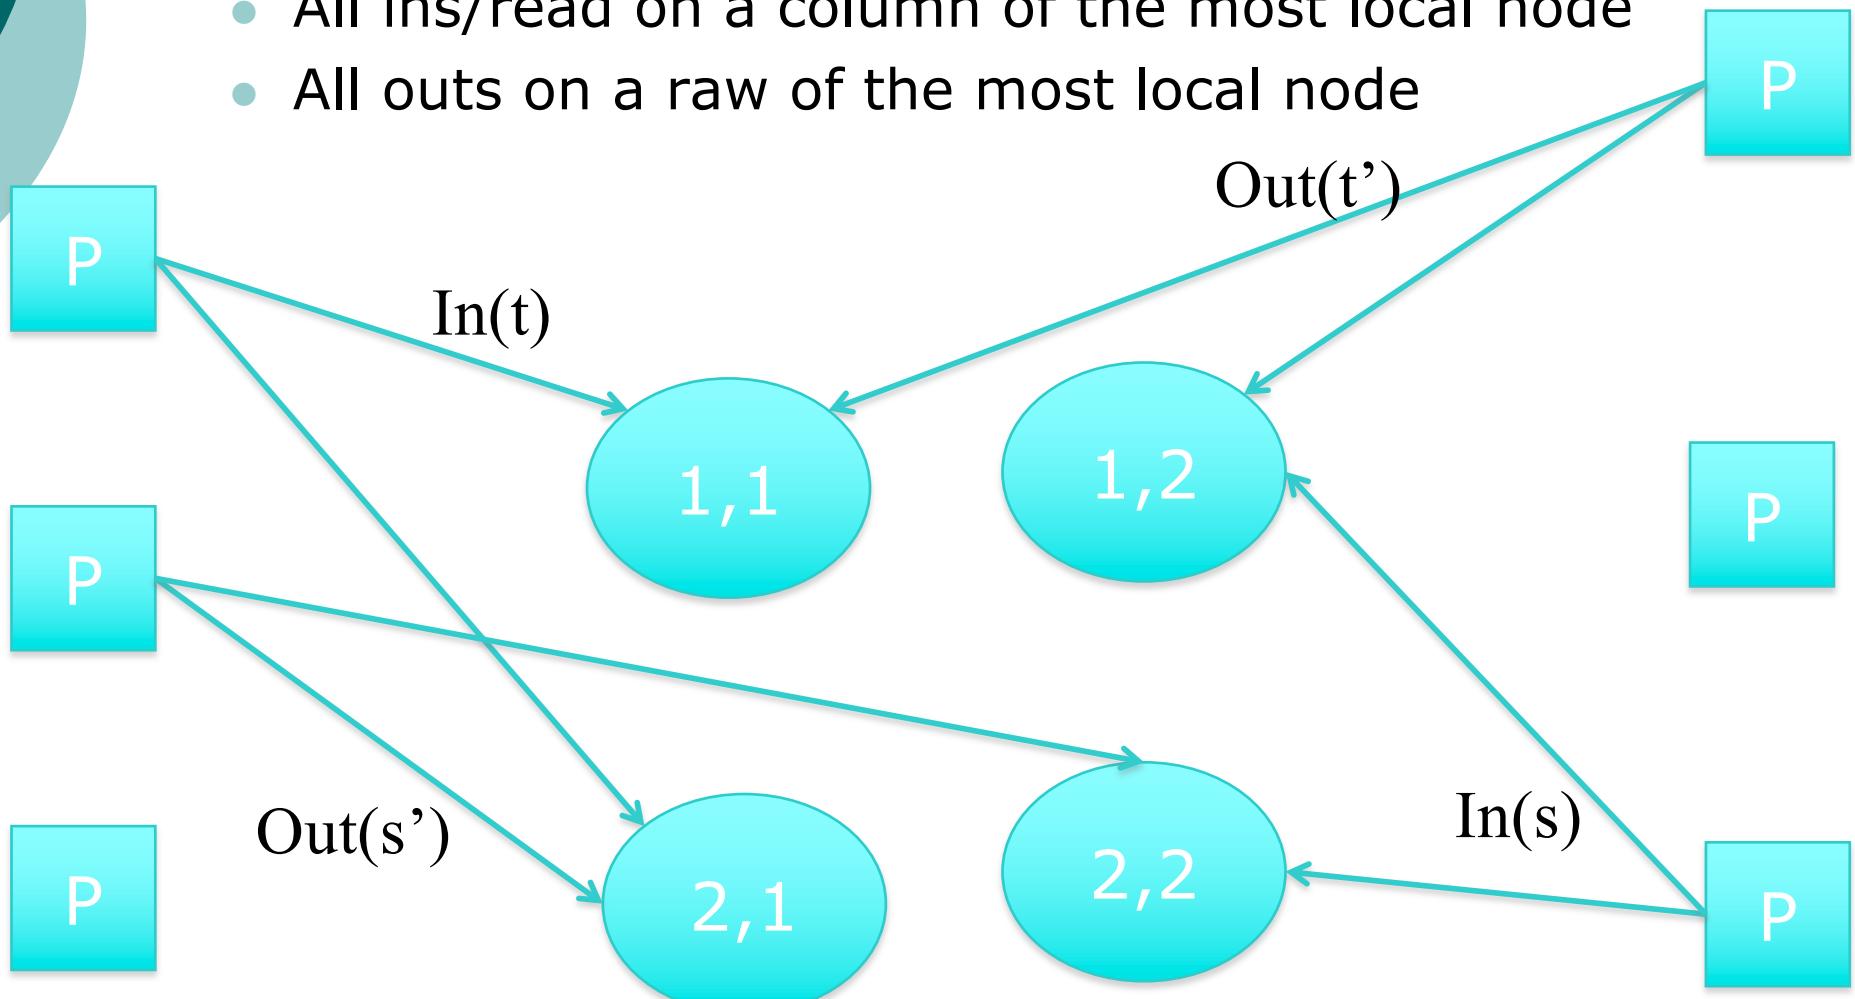
Strategies: Partitionative

- Multicast of outs, Multicast of ins and rds, on a limited number of nodes
- Must ensure matching



Strategies: Matrix Partitionative

- Organize nodes as elements of a virtual matrix
 - All ins/read on a column of the most local node
 - All outs on a raw of the most local node





Strategies: Partitionative

- Reachability
 - Must find a way to ensure it
 - Matrix scheme
 - Hierarchical scheme (F. Zambonelli, 1995); nodes organized as a hierarchy replicating in and out in a whole branch of the hierarchy, from leaves to root
- Consistency
 - Inherits both problems of replicative ins and replicative outs policies
 - But with distributed algorithms that have to execute on a much lower number of nodes
- Pros:
 - Locality in accesses
 - More scalable than fully replicative



Strategies: Considerations

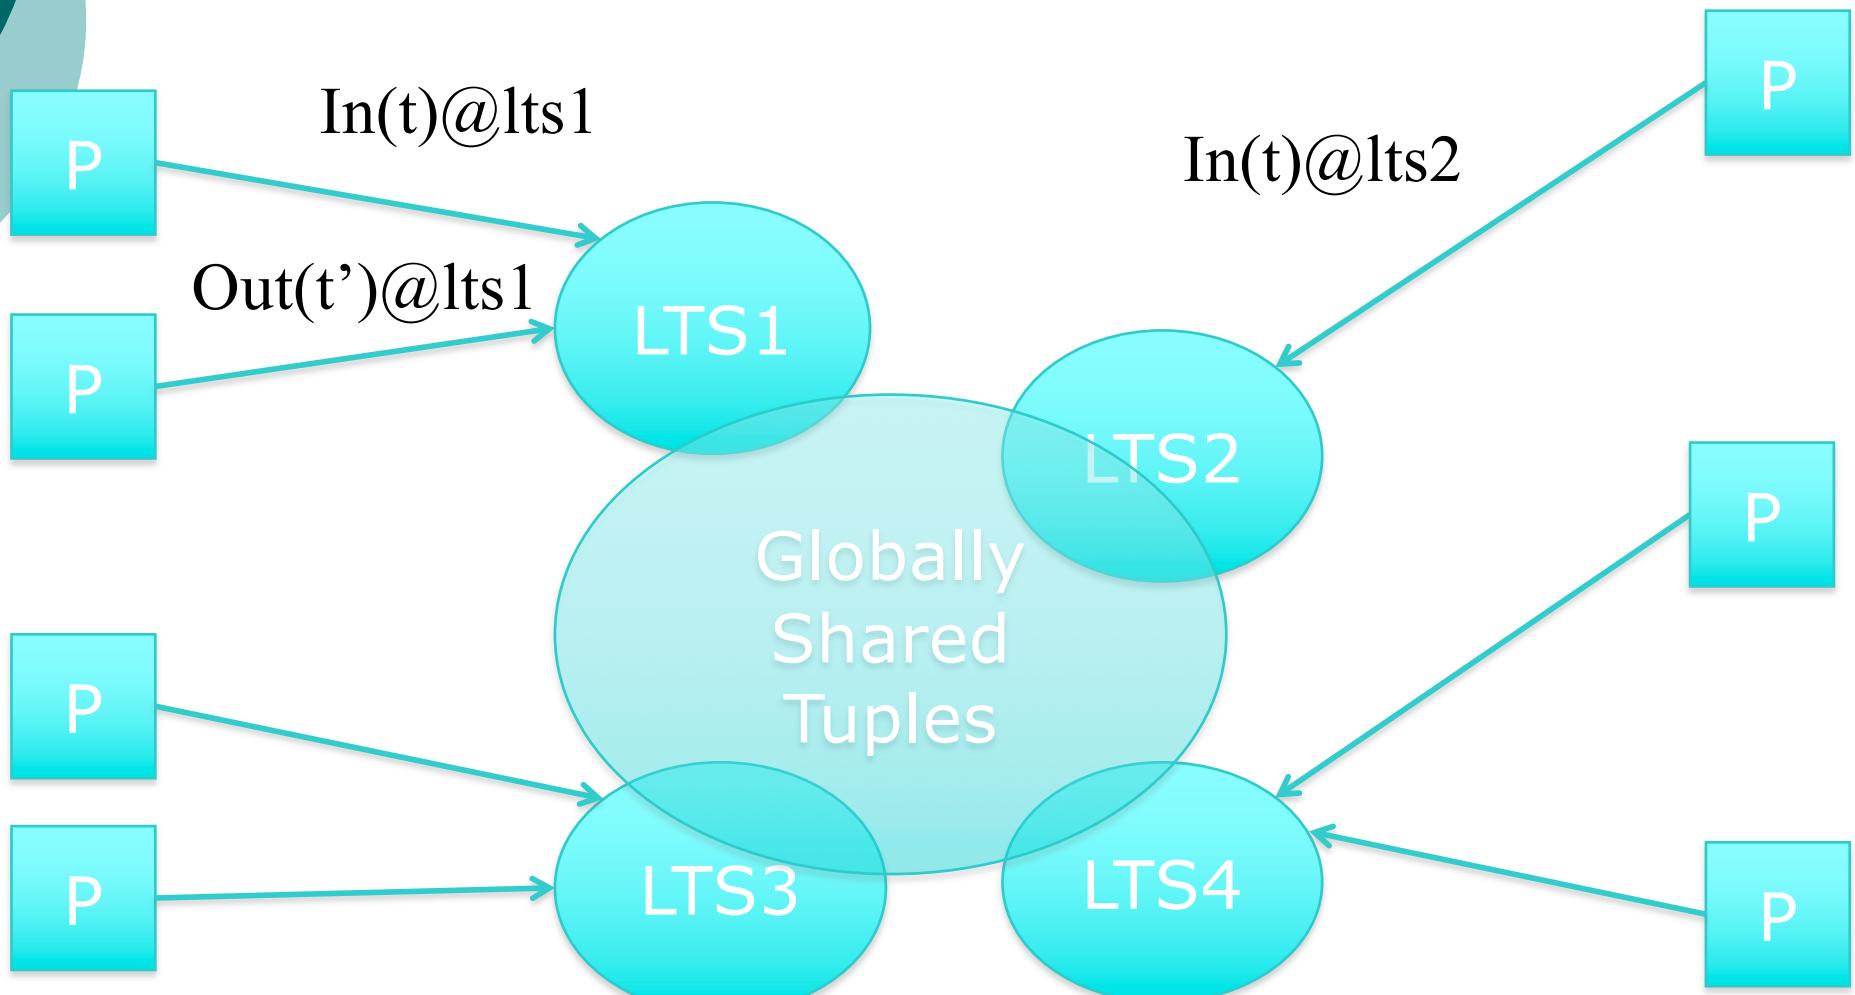
- A very similar problem to that of distributing names in a distributed name space, or events in a distributed event-dispatcher system – Similar strategies have been proposed
- Most existing systems adopts the hash-based scheme or the hierarchical scheme
 - Effectively applicable also to P2P systems, where the “tuples” are the files and each client maps into a portion of the hash function
- **But....is it really always necessary to provide the abstraction of a single global tuple space???**



Local Multiple Tuple Spaces

- Rather than having a single global tuple space, go with
 - A Multiplicity of (mostly) independent tuple spaces
 - Associated to a local logical/physical domain
 - Each with its own local tuple set
 - Possibly sharing only a limited portion of the overall tuples (that is, only a portion of the tuples are “global”)
- Processes
 - Access the local tuple space for most concerns of local computation
 - And sometimes produce “global” tuples
 - And/Or sometimes access other non-local tuples spaces

The General View of Multiple Tuple Spaces





Motivations for Multiple Tuple Spaces (1)

- Separation of Concerns and Encapsulation
 - Encapsulating in a tuple space some data, and giving it access to a limited set of processes
- Security
 - Create access control policies to access to tuple spaces
 - Possibility of giving different grants to different processes in different places



Motivations for Multiple Tuple Spaces (2)

- Inherently “localized” interactions
 - In many applications, processes needs to interact on a location-based way
 - Interact with other processes physically or logically co-located → edge computing
 - Whereas a limited amount of processed needs to interact at a more global scale
- Examples:
 - Intra-organizational processes (logical co-location)
 - IoT - Sensing and acting in the physical world (physical co-location)

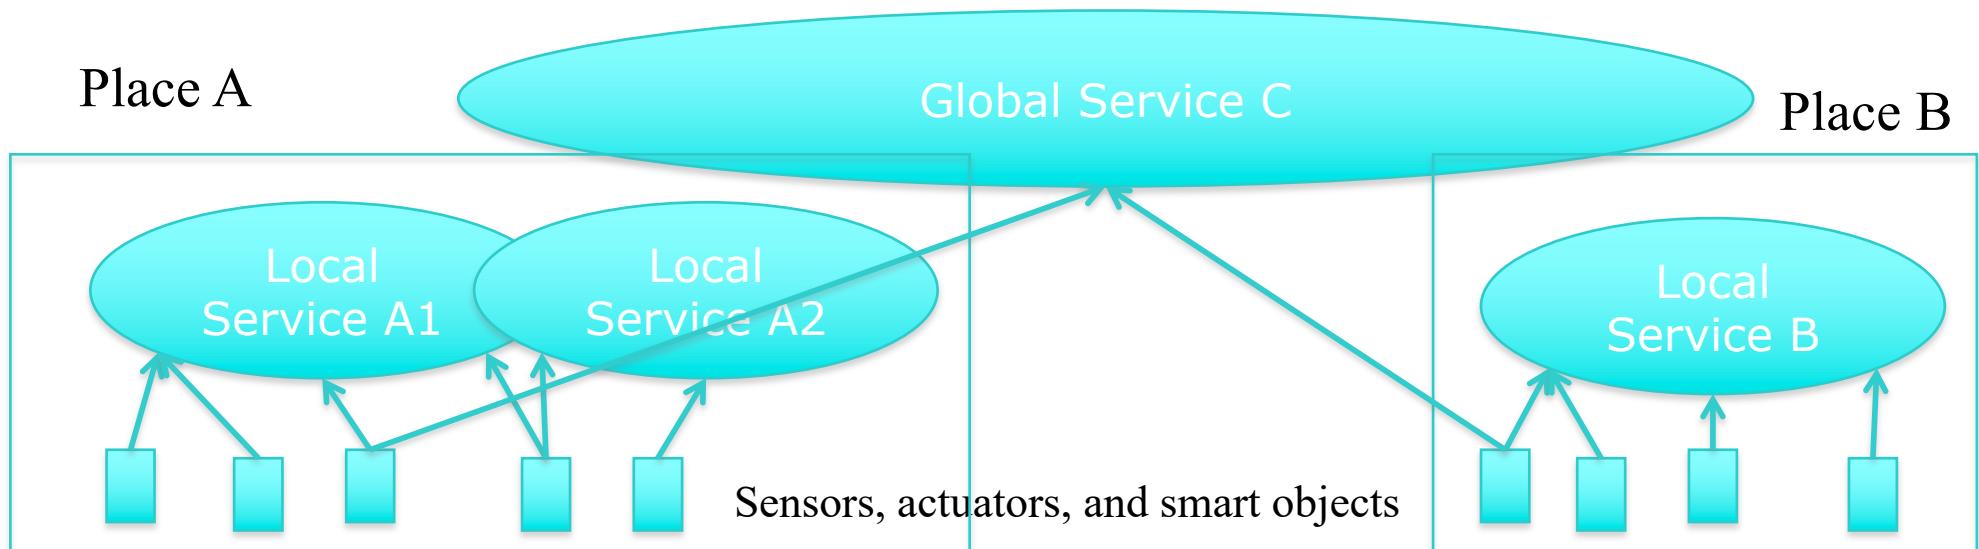


Example: Internet of Things (IoT)

- In the Internet of things vision, Distributed processes associated to sensors and actuators spread in the physical environment and on “smart objects” within → e.g., a smart hotel
- The objects, sensors, and actuators in a room need to interact with each other to provide services, e.g.,
 - Set up windows curtain and lights for lecture presentation in a meeting room
 - Activate sleeping mode in a guest room
 - In the case of detected fire, activate extinguishers, alert via digital signages, open security doors

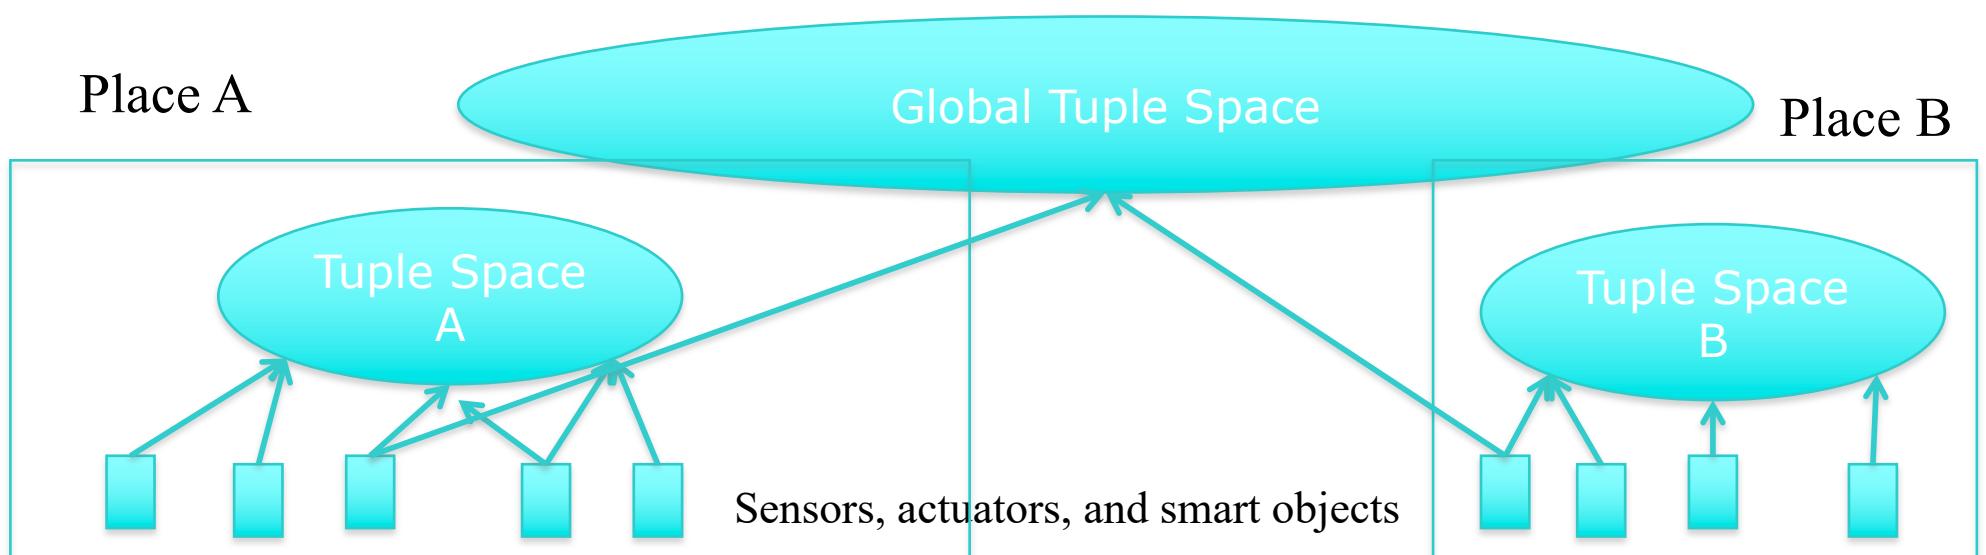
IoT: Local and Global Services

- We can clearly see, in the examples that there are
 - Local services, involving physically localized interactions among local components (e.g., meeting room and guest room services)
 - A few global services, requiring global coordination of some classes of processes at the global level (e.g., fire alarm services)



IoT: Solution with Multiple Local Tuple Spaces

- Local Services are realized by tuple-based coordination of local processes over a local tuple spaces
- Global services are realized (depending on the adopted model) by coordination of global processes
 - By having some of the tuples of interest flow also to the global tuple space, where global processes can access them
 - By having global processes also access local tuple spaces



IoT: Example of Local Service

```
initiate_conference_mode(id)
{
    localTS.out("curtain", "close");
    localTS.out ("lights", "middle_level");
    localTS.out ("projector", "on");
}
```

```
window_curtain()
{ string state, newstate;

    localTS.sub(("curtain", ?new_state), actuate);
}

void actuate(t) // t is the tuple matching the subscription
{if(t.new_state != state)
    actuator.execute(new_state);
}
```



IoT: Example of Global Service (1)

```
temperature_sensor(id)
{ int current_t, alarm_t=60;

    current_t = sensor.getval();

    while(1)
        localTS.out("temp", current_t);
        if (current_t < alarm_t)
            global_TS.out("temp", id, current_t);
}
```

IoT: Example of Global Service (2)

```
digital_signage()
{ string local_message, alarm_message;

    globalTS.sub("alarm", "display", ?alarm_message), actuate);

    localTS.in("display", ?local_message);

        signage.visualize(local_message);
}

void actuate(t) // t is the tuple matching the subscription
{
    signage.visualize(t.alarm_message);
}
```

IoT: Example of Global Service (3)

```
set_up_global_evacuation_plan()
{ int signals = 0, id, int temp;
  global_TS.sub(("temp", ?id, ?temp), start_evacuation);
// don't care who or how many sensor exists
}

void start_evacuation(t)
{ signals++;
  if (signals > 3) // at least 3 sensors have reported

  globalTS.out("alarm", "display", "Escape! doors opening");

  globalTS.out("security_door", "open", "now");
  globalTS.out("extinguisher", "activate", "now");
// don't care who or how many actuator exists
```



Limited Flexibility of Coordination Law: Programmable Laws

- Concept invented by Omicini & Zambonelli, 1998
- A problem of tuple spaces is that the “coordination laws” in the Linda model are always the same
 - basic pattern-matching (defined fields are primitives data types that, if actual, must have the same value)
 - object-matching (e.g., equality of Java serialised form, tuple sub-classing)
- Some times there could be need of more flexible approaches, to avoid complex sequences of ins and outs
 - In(all the sensor tuples with a $20 < \text{temp} < 60$)
 - In (the sensor with the max temp)
- Also, sometimes, tuple spaces associated to different localities may have different rules (also for security purposes)
- How can we do that?



Programmable Coordination Laws

- Programmability of the Coordination language
 - primitives overriding (the default behaviour of a primitive can be changed)
 - adding of new primitives
- Programmability of the Coordination media
 - full monitoring of access events (who, what, on which tuple)
 - which behaviour in response to which access event



Conclusions

- Linda a very effective approach to distributed system programming
 - Simple programming
 - Time and space uncoupling
 - Suited for dynamic scenarios
 - Various implementation possibilities
- Also suited for physically-situated services
 - E.g., sensor networks, Internet of things, robotics



References

- D. Gelernter, N. Carriero, "Coordination Languages and Their Significance", Communications of the ACM, Vol. 35, No. 2, pp. 96-107, February 1992.
- A. Omicini, F. Zambonelli, "Coordination for Internet Application Development", Journal of Autonomous Agents and Multiagent Systems, Vol. 2, No. 3, Sept. 1999.
- G. Cabri, L. Leonardi, F. Zambonelli, "MARS: a Programmable Coordination Architectures for Mobile Agents", IEEE Internet Computing, Vol.5, No. 4, July-August 2000.